

Deploy Generative AI with ease: the OpenVINOTM GenAI library provides developers with necessary tools to optimize and deploy Generative AI models. Based on OpenVINOTM tools and runtime, it provides best performance and characteristics for inference on the supported platforms.

Which use cases are supported by OpenVINOTM GenAI?

OpenVINOTM GenAI provides intuitive and simple C++/Python APIs to run:

- Text generation models (text summarization, rewriting, chatbots, etc.)
- Image generation models (using Diffuser-based architectures)
- Whisper-based speech transcription and translation
- Image processing with Visual Language Models (LLaVa and others)

Use OpenVINOTM GenAI to run models from

 Hugging Face

 ModelScope

Install OpenVINOTM GenAI

Linux install

Windows install

macOS install

[PyPI example](#) for Linux, macOS & Windows:

```
#set up python venv
python -m pip install openvino-genai
```

A full list of installation options is [here](#).

Download converted and optimized models from Hub

Use the huggingface_hub package to download:

```
pip install huggingface_hub
```

Download a model to a local folder:

```
huggingface-cli download "OpenVINO/phi-2-fp16-ov" --local-dir model_path
```

Browse models from OpenVINO Toolkit in our collections:

- [Large Language Models](#)
- [Text2Speech Models](#)

OpenVINOTM GenAI: Supported Models

A list of validated models is [here](#).

Note that even if your model is not listed, it may still work.

Run LLMs for Text Generation



2024 | Updates are [here](#).
Post your questions [here](#).
Read the documentation [here](#).

GenAI Cheat Sheet

1a. Convert and optimize LLMs from Hugging Face



Use the optimum-intel package to convert and optimize models

```
pip install optimum-intel[openvino]
```

Download and convert a model to the OpenVINO IR format while keeping full model precision

```
optimum-cli export openvino --model meta-llama/Llama-2-7b-chat-hf --weight-format fp16 ov_llama_2
```

[Recommended] Download, convert a model, and compress weights to the int4 precision

```
optimum-cli export openvino --model meta-llama/Llama-2-7b-chat-hf --weight-format int4 ov_llama_2
```

A full list of conversion options is [here](#).

Pre-converted LLMs are [here](#)

1b. Convert and optimize LLMs from Model Scope

Use modelscope and optimum-intel packages to convert and optimize models

```
pip install modelscope optimum-intel[openvino]
```

Download the required model to a local folder

```
modelscope download --model 'Qwen/Qwen2-7b' --local_dir model_path
```

[Recommended] convert the model and compress weights to the int4 precision (or int8 or fp16)

```
optimum-cli export openvino -m model_path --task text-generation-with-past --weight-format int4 ov_qwen_25
```

2. Run model using OpenVINO™ GenAI

In Python:

```
import openvino_genai as ov_genai
pipe = ov_genai.LLMPipeline(model_path, "GPU")
print(pipe.generate("What is OpenVINO?", max_length=200))
```

In C++:

```
#include "openvino/genai/llm_pipeline.hpp"
#include <iostream>

int main(int argc, char* argv[]) {
    ov::genai::LLMPipeline pipe(model_path, "GPU");
    std::cout << pipe.generate("What is Large Language Model?",
                                ov::genai::max_new_tokens(200));
}
```

Use CPU or GPU as devices without any other code change

When running LLMs it is also possible to:

- Use different generation parameters (sampling types, etc.)
- Optimize for chat scenarios by using the chat mode
- Load LoRA adapters and dynamically switch between them without recompilation
- Use draft models to accelerate generation via Speculative Decoding

Check out our [Python](#) and [C++](#) samples.

Generate Images using Diffusers



2024 | Updates are [here](#).

Post your questions [here](#).

Read the documentation [here](#).

GenAI Cheat Sheet

1a. Convert and optimize Models from Hugging Face



Use the optimum-intel package to convert and optimize models

```
pip install optimum-intel[openvino]
```

Download and convert a model to the OpenVINO IR format while keeping full model precision

```
optimum-cli export openvino --model stabilityai/stable-diffusion-xl-base-1.0  
--weight-format fp16 ov_SDXL
```

[Recommended] Download, convert the model, and compress weights to the int4 precision

```
optimum-cli export openvino --model stabilityai/stable-diffusion-xl-base-1.0  
--weight-format int4 ov_SDXL
```

A full list of conversion options is [here](#).

1b. Convert and optimize Models from Model Scope



Use modelscope and optimum-intel packages to convert and optimize models

```
pip install modelscope optimum-intel[openvino]
```

Download the required model to a local folder

```
modelscope download --model AI-ModelScope/stable-diffusion-xl-base-1.0 '  
--local_dir model_path
```

[Recommended] convert the model and compress weights to the int4 precision (can use int8 & fp16)

```
optimum-cli export openvino -m model_path --task text-generation-with-past  
--weight-format int4 ov_SDXL
```

2. Run model using OpenVINO™ GenAI

In Python:

```
import openvino_genai as ov_genai  
pipe = openvino_genai.Text2ImagePipeline(model_dir, "GPU")  
image_tensor = pipe.generate(prompt)
```

In C++:

```
#include "openvino/genai/image_generation/text2image_pipeline.hpp"  
  
int main(int argc, char* argv[]) {  
    ov::genai::Text2ImagePipeline pipe(models_path, "GPU");  
    ov::Tensor image = pipe.generate(prompt);  
}
```

Use CPU or GPU as devices without any other code change

When generating images, it is also possible to:

- Alter parameters (width, height, iterations) and compile models for static size
- Load LoRA adapters (in safetensor format) and dynamically switch between them
- Generate multiple images per one request

Check out our [Python](#) and [C++](#) samples.

Processing speech by Whisper



2024 | Updates are [here](#).
Post your questions [here](#).
Read the documentation [here](#).

GenAI Cheat Sheet

1a. Convert and optimize models from Hugging Face



Use the optimum-intel package to convert and optimize models

```
pip install optimum-intel[openvino]
```

Download and convert the model to the OpenVINO format while keeping full model precision

```
optimum-cli export openvino --trust-remote-code --model openai/whisper-base  
ov_whisper
```

Optimization sample is available [here](#).

Pre-converted Models are [here](#)

1b. Convert and optimize models from Model Scope



Use modelscope and optimum-intel packages to convert and optimize models

```
pip install modelscope optimum-intel[openvino]
```

Download the required model to a local folder

```
modelscope download --model 'AI-ModelScope/whisper-large-v3-turbo'  
--local_dir model_path
```

[Recommended] convert the model

```
optimum-cli export openvino -m model_path  
--task automatic-speech-recognition-with-past ov_whisper
```

2. Run model using OpenVINO™ GenAI

In Python:

```
import openvino_genai as ov_genai  
pipe = openvino_genai.WhisperPipeline(model_dir, "CPU")  
#Pipeline expects normalized audio with Sample Rate of 16kHz  
raw_speech = read_wav(...)  
result = pipe.generate(raw_speech)
```

In C++:

```
#include "openvino/genai/whisper_pipeline.hpp"  
  
int main(int argc, char* argv[]) {  
    ov::genai::WhisperPipeline pipe(model_path, "CPU");  
    //Pipeline expects normalized audio with Sample Rate of 16kHz  
    ov::genai::RawSpeechInput raw_speech = read_wav(...);  
    auto result = pipe.generate(raw_speech, config);  
}
```

Use CPU or GPU as devices without any other code change

When running Whisper models, it is also possible to:

- Translate transcription to English
- Predict timestamps
- Process Long-Form (>30 seconds) audio

Check out our [Python](#) and [C++](#) samples.

Analyzing images with VLMs



2024 | Updates are [here](#).
Post your questions [here](#).
Read the documentation [here](#).

GenAI Cheat Sheet

1a. Convert and optimize VLMs from Hugging Face



Use optimum-intel package to convert and optimize models

```
pip install optimum-intel[openvino]
```

Download and convert a model to the OpenVINO IR format while keeping full model precision

```
optimum-cli export openvino --model openbmb/MiniCPM-V-2_6
--trust-remote-code -weight-format fp16 ov_MiniCPM-V-2_6
```

[Recommended] Download, convert the model, and compress weights to the int4 precision

```
optimum-cli export openvino --model openbmb/MiniCPM-V-2_6
--trust-remote-code -weight-format int4 ov_MiniCPM-V-2_6
```

A full list of conversion options is [here](#).

1b. Convert and optimize VLMs from Model Scope



Use modelscope and optimum-intel packages to convert and optimize models

```
pip install modelscope optimum-intel[openvino]
```

Download the required model to a local folder

```
modelscope download --model OpenBMB/MiniCPM-V-2_6 ' --local_dir model_path
```

[Recommended] convert the model and compress weights to the int4 precision (or int8 or fp16)

```
optimum-cli export openvino -m model_path --task image-text-to-text
--weight-format int4 ov_MiniCPM-V-2_6 --trust-remote-code
```

2. Run model using OpenVINO™ GenAI

In Python:

```
import openvino_genai as ov_genai
pipe = ov_genai.VLMPipeline(model_dir, "GPU")
#read images to OV Tensors
rgbImages = read_images(...)
print(pipe.generate(prompt, images=rgbImages, max_new_tokens=100))
```

In C++:

```
#include "openvino/genai/visual_language/pipeline.hpp"
#include <iostream>

int main(int argc, char* argv[]) {
    ov::genai::VLMPipeline pipe(model_dir, "GPU");
    //Read images into vector of OV Tensors
    std::vector<ov::Tensor> rgbImages = read_images(...);
    std::cout << pipe.generate(prompt, ov::genai::images(rgbImages),
                                ov::genai::max_new_tokens(100));
}
```

Use CPU or GPU as devices without any other code change

When running VLMs, it is also possible to:

- Use different generation parameters (sampling types, etc.)
- Optimize for chat scenarios by using chat mode
- Pass multiple images to a model

Check out our [Python](#) and [C++](#) samples.